# News Topic Modeling with Latent Dirichlet Allocation

Alex Jian

April 21, 2023

## 1 Introduction

In this report, I investigate a Bayesian hierarchical model called Latent Dirichlet Allocation (LDA) and use it to build a topic model for a collection of news articles. News articles are written independently, but they generally share common topics and themes - news can be grouped into categories such as politics, health, and finance, and two distinct articles may cover the the same news event from different perspectives or multiple events from the same developing story. Humans reading news articles can easily identify those that cover the same topic, and Latent Dirichlet Allocation allows programs to approximate the same classifications automatically by analyzing the statistical properties of the observed words. This report will both explain how LDA works and apply it to construct a topic model over a large set of news articles.

## 2 Latent Dirichlet Allocation

### 2.1 Overview

Latent Dirichlet Allocation is a generative statistical model which explains a large number of observations through a smaller set of unobserved groups. In the context of machine learning and natural language processing, it is used to discover common topics underlying a collection of documents and then classify each individual document based on which of these topics it contains. In particular, we define a topic as a set of words that collectively suggest a common theme and use LDA to discover these topics based on the repeated appearance of these words. For example, in a news dataset, the words *election, president* or *law* might suggest a politics-related topic, while the words *technology, breakthrough* or *experiment* may suggest a science-related theme. Note that LDA only finds the words that best represent each topic - assigning meaningful labels to these topics requires domain knowledge on the user's part.

## 2.2 Probabilistic Graph Model

As a generative model, Latent Dirichlet Allocation assumes a statistical process by which documents are generated, then conducts Bayesian inference to learn what the parameters of this statistical process are most likely to be. The parameters of the process define the per-document topic and word distributions, which allow for the construction of the topic model over a full document collection.

The probabilistic graph model below shows the generative, topic-based process by which LDA assumes documents are generated. Before showing the diagram, we define variables as follows:

$M$ is the total number of documents in the collection
$N_i$ is total number of words in document $i$
$\alpha$ is a vector specifying the Dirichlet prior for the per-document topic distribution
$\beta$ is a vector specifying the Dirichlet prior for the word distribution per topic
$\theta_i$ is the topic distribution for a single document (document $i$)
$\psi_k$ is the word distribution for a single topic (topic k)
$z_{i_j}$ is the topic referred to by the $jth$ word in document $i$
$w_{i_j}$ is the $jth$ observed word appearing in document $i$



Figure 1:
Graph Model of LDA document generation process

Based on this diagram, we assume that the words for each document in a collection are generated as follows:

1) For each document $i$, choose $\theta_i \sim Dir(\alpha)$. Intuitively, we can think of this as the author deciding which topics the document will cover.

2) For each topic $k$, choose $\psi_k \sim Dir(\beta)$. Intuitively, once the topic is chosen,

certain words are much more likely to appear than others, a property specified by this per-topic word distribution.

3) We assume that each word in each document refers to one singular topic. For each word of each document, first choose its topic, $z_{ij} \sim Categorical(\theta_i)$. Then, choose the word from the topic's word distribution, $w_{i_j} \sim Categorical(\psi_{z_{i_j}})$.

Figure 1 illustrates this process at a high level. Note that W is grayed out, emphasizing that the words are the only observable variable, and the rest are latent variables which must be inferred through Bayesian Inference. Since each node in the graph is conditionally dependent only on its immediate predecessors, we can treat LDA as a Hierarchical Model and conduct inference using Bayesian methods. Note that this model assumes that all words within a document can be treated as exchangable, independent and identically distributed random variables. This assumption is largely untrue - words in sentences depend on each other and have context - but for the purposes of topic modeling, this assumption has still led to good results.

## 2.3   Prior Distributions

In LDA, it is important that the prior distributions accurately reflect our existing knowledge of text documents and their topic and word distributions. In this case, we observe that most documents (especially news articles) are rather focused, and contain only a few topics at most. We can account for this by a constructing sparse symmetric prior ($Dir(\alpha)$, $\alpha < 1$) for the per-document topic distribution, so that while there is no bias against any specific topic, there is bias against more uniform distributions that treat many topics as reasonably likely. The same holds true for the per-topic word distribution ($Dir(\beta)$, $\beta < 1$): for a given topic, certain words are more likely than others, a fact the prior should reflect.



Figure 2:
Sparse Symmetric Prior for 1D Dirichlet Distribution
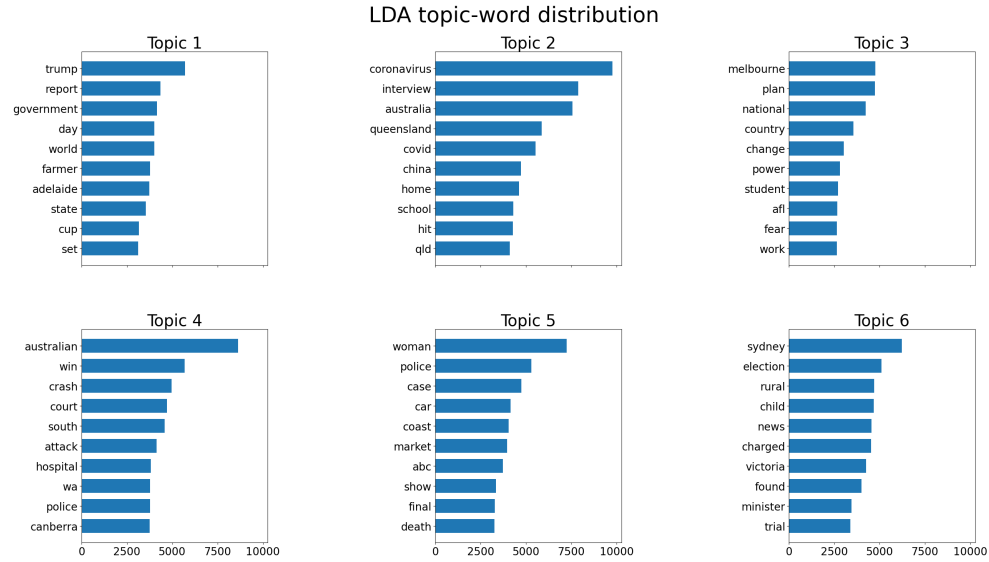
## 2.4   Inference

# 3   Results



Figure 3:
Posterior Topic-Word Distribution



| | publish_date | head_clean | topics |
|---|---|---|---|
| **545241** | 20100620 | call national standard driver licence | Topic 2: 0.594 |
| **618253** | 20110605 | lake face afl mental hurdle eade | Topic 1: 0.417 Topic 2: 0.201 Topic 3: 0.213 |
| **843702** | 20131230 | lebanon buries critic hezbollah | Topic 1: 0.291 Topic 2: 0.266 Topic 3: 0.259 |
| **839535** | 20131206 | glory take point free flowing battle | Topic 2: 0.564 Topic 3: 0.213 |
| **181384** | 20050811 | boy charged sexual assault | Topic 0: 0.227 Topic 4: 0.405 |
| **1225965** | 20201228 | flood rain clermont queensland central quad | Topic 1: 0.490 Topic 5: 0.301 |
| **652951** | 20111116 | dairy industry welcome public committment open | Topic 0: 0.340 |
| **814817** | 20130829 | bundaberg flood money | Topic 3: 0.513 Topic 5: 0.241 |
| **241230** | 20060605 | launceston lead ta population growth | Topic 0: 0.219 Topic 1: 0.233 Topic 3: 0.379 |
| **246294** | 20060629 | union heart big turnout protest | Topic 0: 0.319 Topic 2: 0.261 Topic 3: 0.262 |
| **1203972** | 20200624 | threatened native specie roam central australia | Topic 1: 0.278 Topic 2: 0.208 Topic 4: 0.357 |
| **939127** | 20150319 | india win toss ; bat quarter final bangladesh | Topic 2: 0.214 Topic 3: 0.296 Topic 4: 0.332 |
| **1177263** | 20190925 | ghost net northern gulf hotspot risk marine life | Topic 0: 0.243 Topic 2: 0.231 Topic 5: 0.299 |
| **862232** | 20140324 | arnold palmer invitational live | Topic 1: 0.310 Topic 3: 0.235 Topic 4: 0.270 |

Figure 4:
Posterior per-document Topic Distribution, showing topics which contributed significantly to the
observed words ($> 0.2$)